



# **COMP 2280 - Introduction to Computer Systems**

**Module 3- Introduction to Instruction Set  
Architecture (ISA) and Assembly Language**

# ISA - Introduction

- The instruction set architecture (ISA) provides the interface between computer programs and the hardware.
- The *ISA* is the part of the processor (CPU) that is visible to the programmer.
- Each type of CPU architecture has its own ISA.
- The instruction set of the ISA specifies the commands that can be used to program a CPU or the operations that can be performed by the CPU.

# ISA - Introduction

- You must distinguish between ISA and how it is implemented.
- Example – Consider the IA32 (3<sup>rd</sup> gen of x86 ISA, i386)
  - Intel and AMD have vastly different implementations for the same ISA.
  - Programs written for one will run on the other.
- There are other ISAs, which are not binary compatible with IA32
  - PowerPC (Pre-2005 Macs & Consoles)
  - MIPS
  - Sparc
  - ARM (Mobile devices)
  - M68000 (Motorola 68k)

# ISA - Introduction

- **ISA** = All of the *programmer-visible* components and operations of the computer
  - **memory organization**
    - address space -- how many locations can be addressed?
      - "How many boxes do I have?"
    - addressability -- how many bits per location?
      - "How big are my boxes?"
  - **register set**
    - how many? what size? how are they used?
  - **instruction set**
    - Opcodes (ADD/0001, AND/0101, LEA/1110)
      - Book: The **Opcode** (symbolic name) matches the **opcode** (bits) of the instruction
    - Data types (floating point, integers, etc.)
    - Addressing modes
      - Direct: "Go to place X"
      - Indirect: "Go to the X<sup>th</sup> address of the phonebook"
      - Offset: "Go X steps to the left from your current location"
- ISA provides all information needed for someone that wants to write a program in **machine language** (or translate from a high-level language to machine language).

# ISA - Introduction

- Instructions can be classified as:
  - Data Processing:
    - Arithmetic and Logic instructions
  - Data Storage:
    - Memory instructions
  - Control:
    - Comparison/Test and Branch instructions
  - Data Movement:
    - I/O instructions (not always present – memory mapped I/O)

# ISA-Design

- The number and types of instructions available in the instruction set dictates its complexity:
  - CISC (complex instruction set computer): provides a large collection of instructions, with some support for HLL (high level language). (x86/x64 is CISC-ish)
  - RISC (reduced instruction set computer): provides fewer instructions, and usually no support for HLL. Idea behind RISC is:
    - Make the common case fast
    - 80:20 rule (use freed space for efficient implementations)
    - ARM is RISC
      - Previously **Advanced RISC Machine**, originally **Acorn RISC Machine**

# ISA Design

- There are tradeoffs between a small and large instruction set.
  - Large instruction sets requires more logic to implement, but is easier to use.
  - Small instruction sets are easier to implement, but harder to use.

# ISA Design

- Examples of RISC Architectures
  - Sparc, MIPS, IBM RS/6000
- Examples of CISC Architectures
  - Motorola 68000, Intel 80386

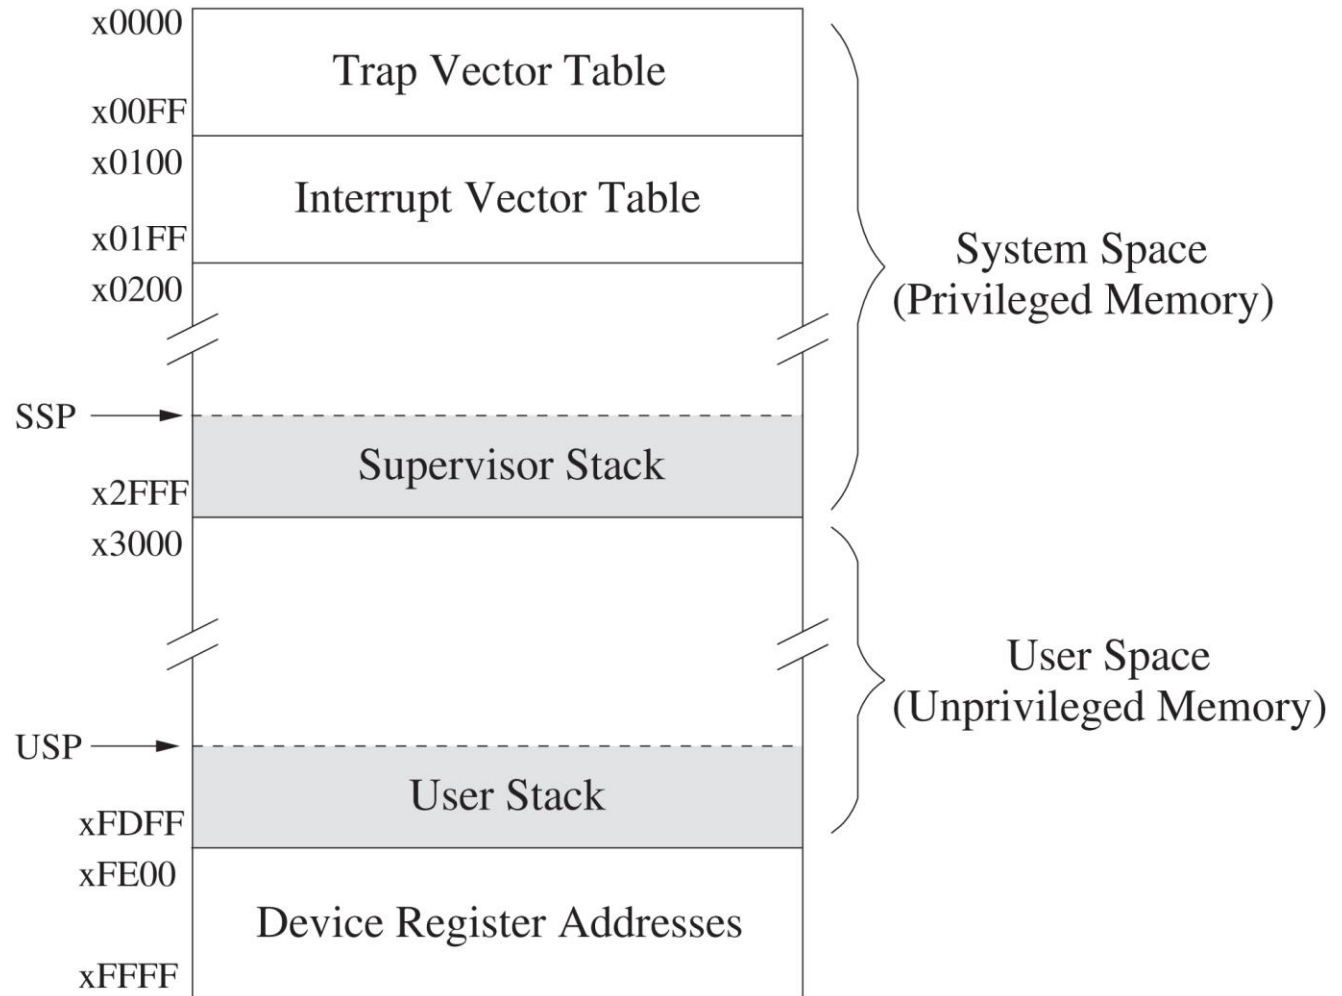


# LC3 ISA

- Let us take a look at the LC3 ISA.
  - *Appendix A* of the textbook.
- Memory
  - address space:  $2^{16}$  locations (16-bit addresses)
  - addressability: 16 bits ie. Memory is addressed 16 bits at a time.
  - Here is a layout of LC-3's memory (memory map)

# LC3 ISA

Copyright © McGraw-Hill Education. Permission required for reproduction or display.



Memory Map of the LC-3

# LC3 ISA

- Registers

- *temporary high-speed storage, accessed in a single machine cycle*
  - *accessing regular memory (RAM) generally takes longer than a single cycle*
  - *“Sort of / kind of” like a variable but is not a variable*
- **eight general-purpose registers: R0 - R7**
  - each 16 bits wide
- **other registers**
  - not directly addressable, but used by (and affected by) instructions
  - PC (program counter)
  - IR (instruction register)
  - condition codes (N = Negative, Z = Zero, P = Positive)

# LC3 ISA

- Opcodes (actual instructions available)
  - 15 opcodes
    - Each instruction is 16-bits long, with 4 bit opcode
    - 12 bits for operands and addressing mode
  - *Operate* instructions: ADD, AND, NOT
  - *Data movement* instructions: LD, LDI, LDR, LEA, ST, STR, STI
  - *Control* instructions: BR, JSR/JSRR, JMP/RET, RTI, TRAP
    - some opcodes set/clear *condition codes*, based on result:
      - N = negative, Z = zero, P = positive ( $> 0$ )

# LC3 ISA

- Data Types
  - 16-bit 2's complement integer
- Addressing Modes
  - How is the location of an operand specified?
  - non-memory addresses: *immediate*, *register*
  - *memory* addresses: *PC-relative*, *indirect*, *base+offset*

# LC3 ISA

- Only three arithmetic/logic instructions: ADD, AND, NOT
  - Source and destination operands are registers
    - These instructions do not reference memory.
    - ADD and AND can use “immediate” mode, where one operand is hard-wired into the instruction.

# Human-Friendly Programming

- Computers need binary instruction encodings...
  - 0001110010000110
- Humans prefer symbolic languages...
  - $a = b + c$
- High-level languages allow us to write programs in clear, precise language that is more like English or math. Requires a program (compiler) to translage from symbolic language to machine instructions.


# ISA Programming

- Assembly Language is a low-level symbolic language, just a short step above machine instructions.
- Don't have to remember opcodes (ADD = 0001, NOT = 1001, ...).
- Give symbolic names to memory locations -- don't have to do binary arithmetic to calculate offsets.
- Like machine instructions, allows programmer explicit, instruction-level specification of program.
- Disadvantage:
  - Not portable. Every ISA has its own assembly language. Program written for one platform does not run on another.



# Assembly Language

- Very similar format to instructions -- replace bit fields with symbols.



- For the most part, one line of assembly language = one instruction.
- Some additional features for allocating memory, initializing memory locations, service calls.
- Numerical values specified in hexadecimal (x30AB) or decimal (#10).

# Example Program

```
;
; Program to multiply a number by the constant 6
;
        .ORIG    x3050
        LD       R1, SIX
        LD       R2, NUMBER
        AND      R3, R3, #0           ; Clear R3.  It will
                                       ; contain the product.
; The inner loop
;
AGAIN    ADD      R3, R3, R2
        ADD      R1, R1, #-1         ; R1 keeps track of
        BRp      AGAIN              ; the iteration.
;
        HALT
;
NUMBER   .BLKW    1
SIX      .FILL    x0006
;
        .END
```

Instructions

Comments

Assembler Directives

Labels

# Assembly Language

- **LC-3 Assembly Language Syntax**
  - Each line of a program is one of the following:
    - an instruction
    - an assembler directive (or pseudo-op)
    - a comment
  - White space (between symbols) and case is ignored.
  - Comments (beginning with “;”) are also ignored.
    - An instruction has the following format:  
**LABEL OPCODE OPERANDS ; COMMENTS**
      - Bold indicates required, others are optional

# Assembly Language

- **Mnemonics**

- reserved symbols that correspond to LC-3 instructions
- listed in Appendix A
  - ex: ADD, AND, LD, LDR, ...

- **Operands**

- registers -- specified by Rn, where n is the register number
- numbers -- indicated by # (decimal) or x (hex)
- label -- symbolic name of memory location
- separated by comma
- number, order, and type correspond to instruction format
- ex:

ADD	R1,R1,R3
ADD	R1,R1,#3
LD	R6,NUMBER
BRz	LOOP

# Assembly Language

- **Label**

- placed at the beginning of the line
- assigns a symbolic name to the address of the corresponding line
- “Sort of / kind of” like a pointer, but is not a pointer
  - Ex: “*Alex’s House*” is the label for the address “22 *Something St*”
- ex:  

LOOP	ADD	RI,RI,#-1
	BRp	LOOP

- **Comment**

- anything after a semicolon is a comment
- ignored by assembler
- used by humans to document/understand programs
- tips for useful comments:
  - avoid restating the obvious, as “decrement RI”
    - What is RI?
  - provide additional insight, as in “accumulate product in R6”
  - use comments to separate pieces of program
  - Use comments to make the graders lives easier
    - Happy grader, happy grading, happy gradee

# Assembler Directives

- **Pseudo-operations (7.2.2 p.236)**

- do not refer to operations executed by program
- used by assembler
- look like instructions, but “mnemonic” starts with a dot (.)
- Eg)

```
data        .fill  #1000 ;1000 dec
sum         .fill  x500  ;500 hex
```

# Assembler Directives

<i>Directive</i>	<i>Operand</i>	<i>Meaning</i>
.ORIG	address	starting address of program
.END		end of program
.BLKW	n	allocate n blank words (aka: memory addresses) of storage
.FILL	n	allocate one word, initialize with value n
.STRINGZ	n-character string	allocate n+1 locations, initialize w/characters and null terminator

# Traps

- LC-3 assembler provides “pseudo-instructions” for each trap code, so you don’t have to remember them.
- A trap provides a mechanism for user (unprivileged) code to execute OS (privileged) code.
- “Pretty much” a system call
- Useful in any system that has an OS.
- For us, traps provide a way to do I/O and halt the machine.



# Traps


<i>Name</i>	<i>Equivalent</i>	<i>Description</i>
HALT	TRAP x25	Halt execution and print message to console.
IN	TRAP x23	Print prompt on console, read (and echo) one character from keybd. Character stored in R0[7:0].
OUT	TRAP x21	Write one character (in R0[7:0]) to console.
GETC	TRAP x20	Read one character from keyboard. Character stored in R0[7:0].
PUTS	TRAP x22	Write null-terminated string to console. Address of string is in R0.

# Assembly Process

- **The Assembly Process (7.3)**
  - Objective
    - Translate the AL (Assembly Language) program into ML (Machine Language).
    - Each AL instruction yields one ML instruction word.
    - Interpret pseudo-ops correctly.
  - Problem
    - An instruction may reference a label.
    - If the label hasn't been defined yet, the assembler can't form the instruction word
  - Solution
    - Two-pass assembly

# Assembly Process

- Convert assembly language file (.asm) into an executable file (.obj) for the LC-3 simulator.



- **First Pass:**
  - scan program file
  - find all labels and calculate the corresponding addresses; this is called the symbol table
- **Second Pass:**
  - convert instructions to machine language, using information from symbol table

# Assembly Process

- **First Pass - generating the symbol table**
- 1. Find the .ORIG statement, which tells us the address of the first instruction.
  - Initialize location counter (LC), which keeps track of the location of the current instruction.
- 2. For each non-empty line in the program:
  - a) If line defines a label, add label and LC to symbol table.
  - b) Increment LC.
    - NOTE: If statement is .BLKW or .STRINGZ, increment LC by the number of words allocated.
- 3. Stop when .END statement is reached.

# Assembly Process

```
1  .ORIG    x3050
2          LD   R1, SIX
3          LD   R2, NUMBER
4          AND  R3, R3, #0
5  ;
6  ; The inner loop
7  ;
8  AGAIN
9          ADD  R3, R3, R2
10         ADD  R1, R1, b-01
11         BRp  AGAIN
12        ;
13        HALT
14
15  NUMBER  .BLKW #2
16         .STRINGZ "Orphan Text"
17  SIX     .FILL  x0006
18
19  .end
```

Symbol	Address
Again	
Number	
Six	

## Obscure sidenote:

- Along with '#' for decimal and 'x' for hexadecimal, binary numbers can also be written in LC3 using 'b'
- This course does not use it in any way, shape, or form other than this example.
- Unless specified otherwise, please use '#' or 'x' and **do not use 'b'** in your assignments or labs

# Assembly Process

```
;
; Program to multiply a number by six
;

        .ORIG      x3050
x3050    LD          R1, SIX
x3051    LD          R2, NUMBER
x3052    AND         R3, R3, #0
;
; The inner loop
;
x3053    AGAIN      ADD         R3, R3, R2
x3054    ADD         R1, R1, b-01
x3055    BRp        AGAIN
;
x3056    HALT
;
x3057    NUMBER     .BLKW      #2
x3058    .STRINGZ   "Orphan Text"
x3064    SIX        .FILL      x0006
;

        .END
```

Symbol	Address
Again	x3053
Number	x3057
Six	x3065

## Sidenote:

- We have no way to (directly) reference this string...
- Can you think of a work-around without adding a label?

# Assembly Process

Memory				
❗ ▶	<b>x3050</b>	x2214 8724	LD R1, SIX	
❗ ▶	<b>x3051</b>	x2405 9221	LD R2, NUMBER	
❗ ▶	<b>x3052</b>	x56E0 22240	AND R3, R3, #0	
❗ ▶	<b>x3053</b>	x16C2 5826	ADD R3, R3, R2	
❗ ▶	<b>x3054</b>	x127F 4735	ADD R1, R1, b-01	
❗ ▶	<b>x3055</b>	x03FD 1021	BRp AGAIN	
❗ ▶	<b>x3056</b>	xF025 61477	HALT	
❗ ▶	<b>x3057</b>	x0000 0	NUMBER .BLKW #2	
❗ ▶	<b>x3058</b>	x0000 0	NUMBER .BLKW #2	
❗ ▶	<b>x3059</b>	x004F 79	O	
❗ ▶	<b>x305A</b>	x0072 114	r	
❗ ▶	<b>x305B</b>	x0070 112	p	
❗ ▶	<b>x305C</b>	x0068 104	h	
❗ ▶	<b>x305D</b>	x0061 97	a	
❗ ▶	<b>x305E</b>	x006E 110	n	
❗ ▶	<b>x305F</b>	x0020 32		
❗ ▶	<b>x3060</b>	x0054 84	T	
❗ ▶	<b>x3061</b>	x0065 101	e	
❗ ▶	<b>x3062</b>	x0078 120	x	
❗ ▶	<b>x3063</b>	x0074 116	t	
❗ ▶	<b>x3064</b>	x0000 0	.STRINGZ "Orphan Text"	
❗ ▶	<b>x3065</b>	x0006 6	SIX .FILL x0006	

## Sidenote:

- LC3 shows the memory address, the value at that address (hex and decimal) and the associated instruction line if there is one

Symbol	Address
Again	x3053
Number	x3057
Six	x3065

## Assembled code:

- Note the instruction layout (opcode, operands and addressing modes)

## Sidenote:

- Like in C, strings end in a null character (x0)
- LC3 places the instruction line there because it's the only spot that makes sense since all strings have a null character

# Assembly Process

- **Second Pass - generating the ML program**
  - Scan each line again
  - Translate each AL instruction into ML
    - *Look up symbols in the symbol table*
    - *Ensure that labels are no more than +256 / -255 words from instruction*
    - *Determine operand field for the instruction*
  - Fill memory locations as directed by pseudo-ops
  - Stop when .END is encountered
  - Potential problems:
    - Improper number or type of arguments
      - ex:       NOT       R1,#7  
          ADD       R1,R2  
          ADD       R3,R3,NUMBER
    - Immediate argument too large
      - ex:       ADD       R1,R2,#1023
    - Address (associated with label) more than +256/-255 words from instruction
      - can't use PC-relative addressing mode



# Object File Format

- LC-3 object file contains
  - Starting address (location where program must be loaded), followed by...
  - Machine instructions

```
0011000000000000
```

← .ORIG x3000

```
0101010010100000
```

← AND R2, R2, #0

```
0010011000010001
```

← LD R3, PTR

```
1111000000100011
```

← TRAP x23

```
.
```

```
.
```

```
.
```

# Beyond a Single Object File

- Larger programs may be written by multiple programmers, or may use modules written by a third party. Each module is assembled independently, each creating its own **object file** and **symbol table**.
- To execute, a program must have all of its modules combined into a single **executable** image.
- **Linking** is the process to combine all of the necessary object files into a single executable.

# External Symbols

**Note: LC3 Instruction  
set reference: pages  
655-674, 3rd edition**

- In the assembly code we're writing, we may want to symbolically refer to information defined in a different module.
- For example, suppose we don't know the starting address of the file in our counting program. The starting address and the file data could be defined in a different module.
- We want to do this:
  - `PTR .FILL STARTofFILE`
- To tell the assembler that `STARTofFILE` will be defined in a different module, we could do something like this:
  - `.EXTERNAL STARTofFILE`
- This tells the assembler that it's not an error that `STARTofFILE` is not defined. It will be up to the linker to find the symbol in a different module and fill in the information when creating the executable.